

WileyWinters_Week6_Assignment

February 24, 2024

0.0.1 Week 6 Assignment

- Wiley Winters
 - MSDS 670 — Data Visualization
 - 25-FEB-2024
-

0.0.2 Dataset Information

Dataset: Jobs and Salaries in Data Science Metadata: - **work_year**: Year in which data was recorded. - **job_title**: Specific title of the job role. - **job_category**: Classification of the job role into broader categories for easier analysis - **salary_currency**: Currency in which the salary is paid - **salary**: Annual gross salary of the role in the local currency - **salary_in_usd**: Annual gross salary in USD - **employee_residence**: Country of residence - **experience_level**: Classifies the professional experience level of the employee - **employment_type**: Specifies the type of employment such as *full-time*, *part-time*, *contract*, *etc* - **work_setting**: Work setting or environment such as *remote*, *in-person*, or *hybrid* - **company_location**: Country where the company is located - **company_size**: Size of the employer company categorized as *small (S)*, *medium (M)*, and *large (L)*

Formal Reference to Dataset

Qaasim, H. (2023, December). Jobs and Salaries in Data Science. Version 6. Retrieved December 25, 2023 from <https://www.kaggle.com/datasets/hummaamqaasim/jobs-in-data/data>

Import required packages and libraries. Set global configuration items.

```
[1]: import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
import matplotlib.ticker as mtick
from matplotlib import rcParams
import numpy as np

# Suppress Warnings
import warnings
warnings.filterwarnings('ignore')

# Set matplotliblab autolayout to True
```

```
rcParams.update({'figure.autolayout': True})
```

Read dataset into a Pandas DataFrame

```
[2]: jobs_df = pd.read_csv('../data/jobs_in_data.csv')
jobs_df.sample(5)
```

```
[2]:
```

	work_year	job_title	job_category	\
7517	2022	Data Analyst	Data Analysis	
5348	2023	Data Manager	Leadership and Management	
9157	2020	Research Scientist	Data Science and Research	
7581	2022	Data Engineer	Data Engineering	
6954	2023	Data Architect	Data Architecture and Modeling	

	salary_currency	salary	salary_in_usd	employee_residence	\
7517	USD	150000	150000	United States	
5348	USD	110000	110000	United States	
9157	USD	450000	450000	United States	
7581	USD	75000	75000	United States	
6954	USD	115000	115000	United States	

	experience_level	employment_type	work_setting	company_location	\
7517	Mid-level	Full-time	In-person	United States	
5348	Mid-level	Full-time	In-person	United States	
9157	Mid-level	Full-time	In-person	United States	
7581	Senior	Full-time	In-person	United States	
6954	Senior	Full-time	Remote	United States	

	company_size
7517	M
5348	M
9157	M
7581	M
6954	M

```
[3]: jobs_df.describe().T
```

```
[3]:
```

	count	mean	std	min	25%	\
work_year	9355.0	2022.760449	0.519470	2020.0	2023.0	
salary	9355.0	149927.981293	63608.835387	14000.0	105200.0	
salary_in_usd	9355.0	150299.495564	63177.372024	15000.0	105700.0	

	50%	75%	max
work_year	2023.0	2023.0	2023.0
salary	143860.0	187000.0	450000.0
salary_in_usd	143000.0	186723.0	450000.0

The dataset covers years from 2020 to 2023. In order to not double count some values. I will only

work with 2023 data

Check some basic items to see if the dataset requires cleaning or not

```
[4]: print(jobs_df.info())
print('\nNaN Values:\n', jobs_df.isna().sum())
print('\nDuplicates: ', jobs_df.duplicated().sum())
print('\nSize: ', jobs_df.size)
print('\nDistribution:\n', jobs_df.describe().T)
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 9355 entries, 0 to 9354
Data columns (total 12 columns):
#   Column                Non-Null Count  Dtype
---  -
0   work_year              9355 non-null   int64
1   job_title              9355 non-null   object
2   job_category           9355 non-null   object
3   salary_currency        9355 non-null   object
4   salary                 9355 non-null   int64
5   salary_in_usd          9355 non-null   int64
6   employee_residence     9355 non-null   object
7   experience_level        9355 non-null   object
8   employment_type        9355 non-null   object
9   work_setting           9355 non-null   object
10  company_location       9355 non-null   object
11  company_size           9355 non-null   object
dtypes: int64(3), object(9)
memory usage: 877.2+ KB
None
```

```
NaN Values:
work_year          0
job_title          0
job_category       0
salary_currency    0
salary            0
salary_in_usd     0
employee_residence 0
experience_level   0
employment_type    0
work_setting       0
company_location   0
company_size       0
dtype: int64
```

```
Duplicates: 4014
```

```
Size: 112260
```

Distribution:

	count	mean	std	min	25%	\
work_year	9355.0	2022.760449	0.519470	2020.0	2023.0	
salary	9355.0	149927.981293	63608.835387	14000.0	105200.0	
salary_in_usd	9355.0	150299.495564	63177.372024	15000.0	105700.0	

	50%	75%	max
work_year	2023.0	2023.0	2023.0
salary	143860.0	187000.0	450000.0
salary_in_usd	143000.0	186723.0	450000.0

Looks like there is a lot of duplicates. I will remove them.

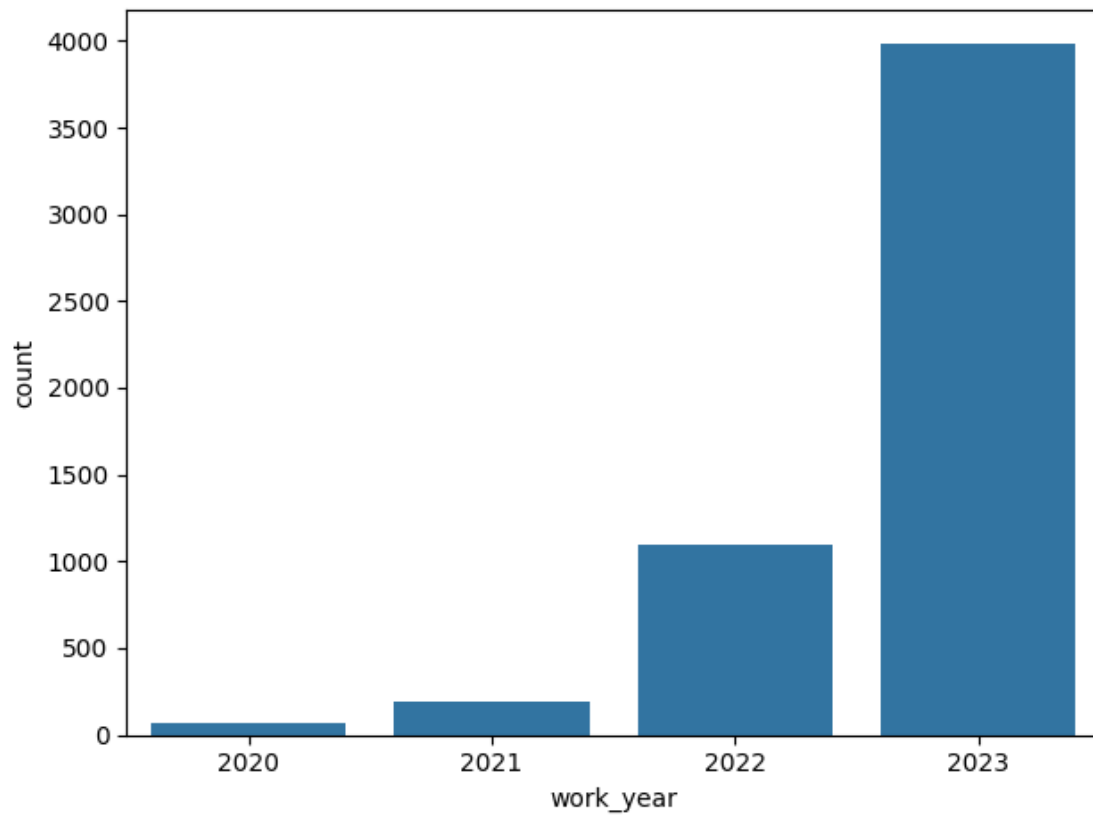
```
[5]: jobs_df.drop_duplicates(keep='first', inplace=True)
jobs_df.duplicated().sum()
```

```
[5]: 0
```

0.0.3 Basic EDA

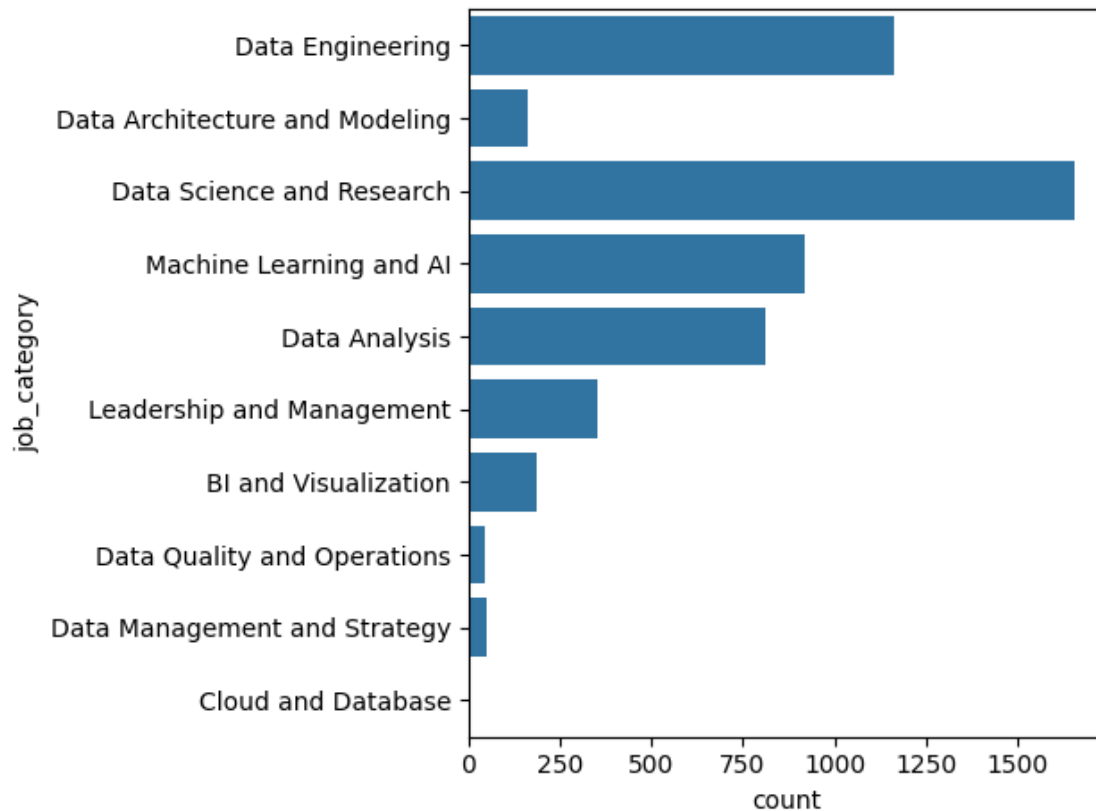
```
[6]: sns.countplot(jobs_df, x='work_year')
```

```
[6]: <Axes: xlabel='work_year', ylabel='count'>
```



```
[7]: sns.countplot(jobs_df, y='job_category')
```

```
[7]: <Axes: xlabel='count', ylabel='job_category'>
```

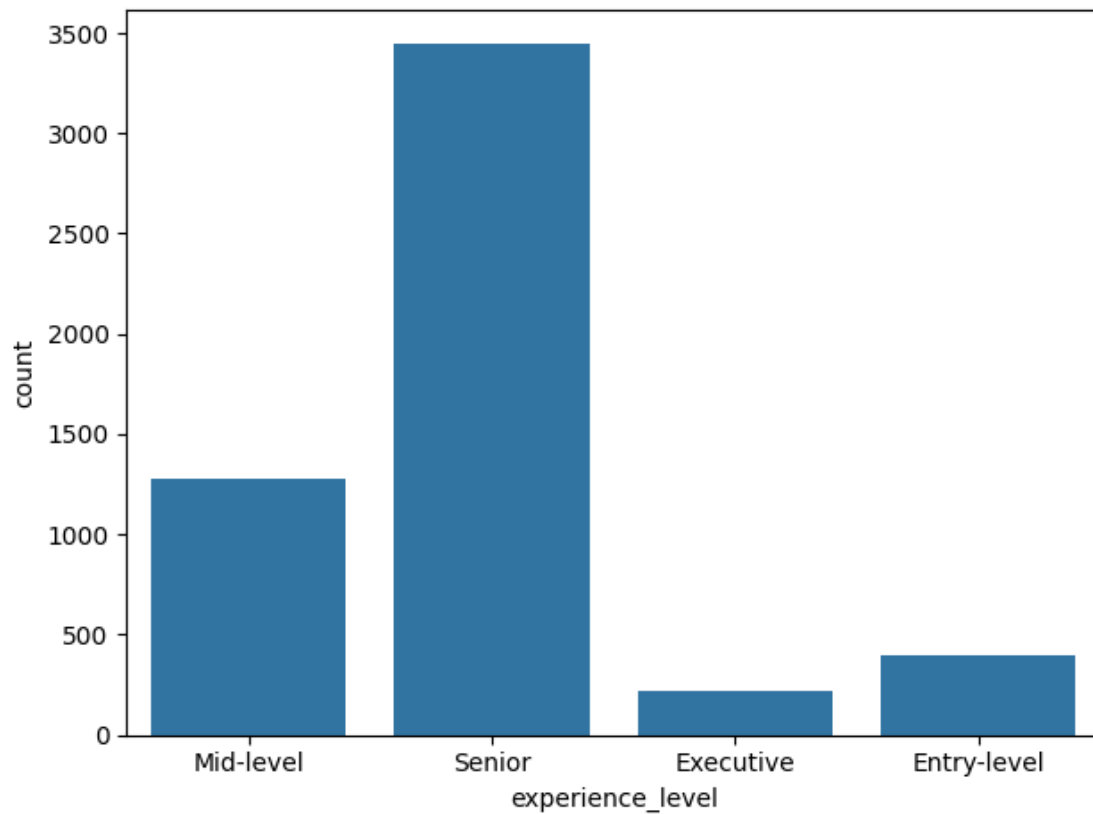


```
[8]: jobs_df['employee_residence'].value_counts().head(10)
#sns.countplot(data=jobs_df, y='employee_residence')
```

```
[8]: employee_residence
United States      4255
United Kingdom     351
Canada             196
Germany            65
Spain              63
France             53
Portugal           26
Netherlands        21
Italy              20
Brazil             19
Name: count, dtype: int64
```

```
[9]: sns.countplot(data=jobs_df, x='experience_level')
```

```
[9]: <Axes: xlabel='experience_level', ylabel='count'>
```




```
[10]: jobs_df['experience_level'].value_counts()
```

```
[10]: experience_level
Senior      3444
Mid-level   1274
Entry-level  400
Executive    223
Name: count, dtype: int64
```

```
[11]: sns.countplot(data=jobs_df, x='employment_type')
```

```
[11]: <Axes: xlabel='employment_type', ylabel='count'>
```

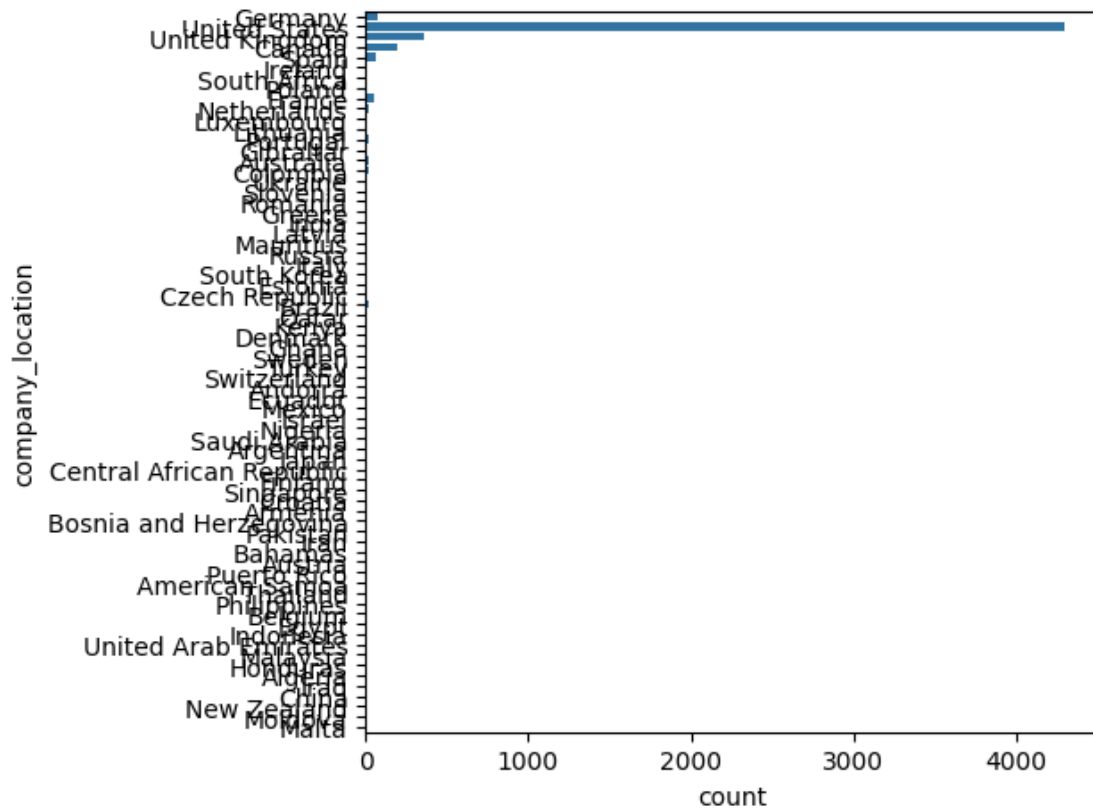


```
[12]: jobs_df['employment_type'].value_counts()
```

```
[12]: employment_type
Full-time    5296
Contract      19
Part-time    15
Freelance     11
Name: count, dtype: int64
```

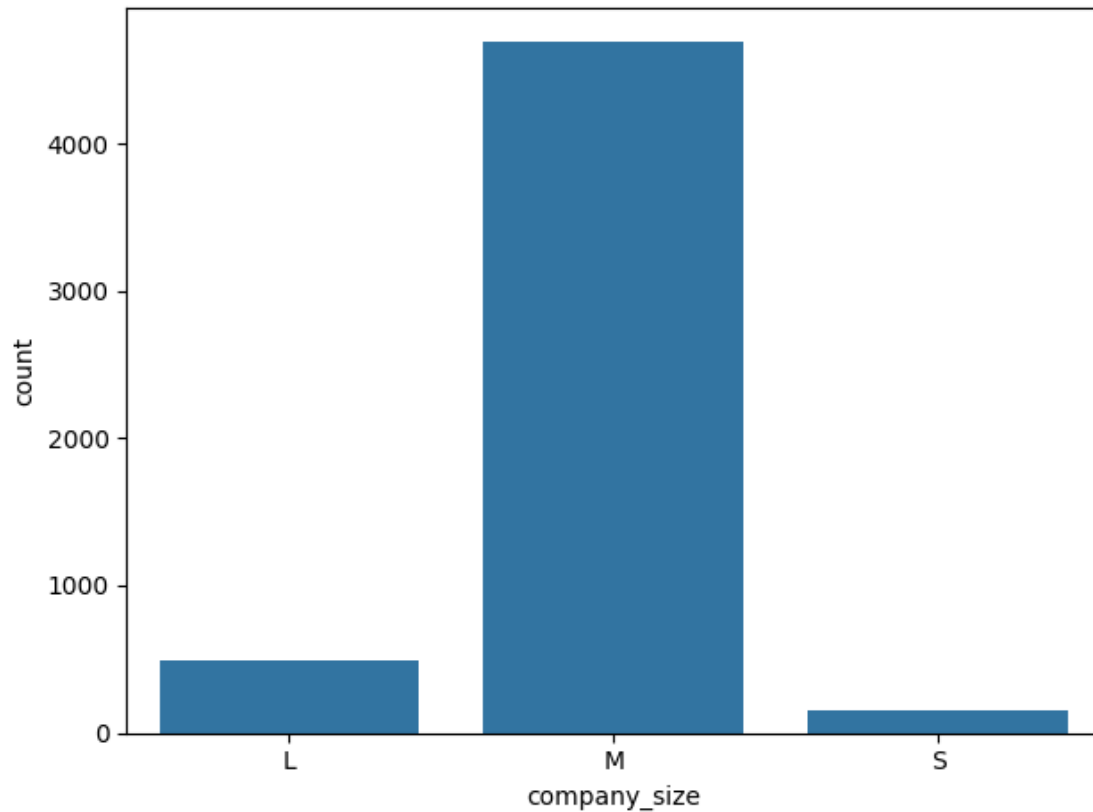
```
[13]: sns.countplot(data=jobs_df, y='company_location')
```

```
[13]: <Axes: xlabel='count', ylabel='company_location'>
```

```
[14]: sns.countplot(data=jobs_df, x='company_size')
```

```
[14]: <Axes: xlabel='company_size', ylabel='count'>
```



Based on EDA performed, I will limit the number of `company_location` and `employee_residence` to 10 countries. I found that some of the countries are only listed a couple of time; therefore, are not well distributed.

```
[15]: jobs_df.value_counts(['company_location', 'employee_residence']).head(15)
```

```
[15]: company_location  employee_residence
United States         United States         4249
United Kingdom        United Kingdom         350
Canada                Canada                 192
Germany               Germany                 60
Spain                 Spain                   57
France                France                   46
Portugal              Portugal                 23
Netherlands           Netherlands              19
Brazil                Brazil                   17
Australia             Australia                17
Colombia              Colombia                 14
Italy                 Italy                   13
Greece                Greece                   11
Mexico                Mexico                    9
```

```
Ireland      Ireland      8
Name: count, dtype: int64
```

```
[16]: countries = ['United States', 'United Kingdom', 'Canada', 'Germany',
                  'Spain', 'France', 'Portugal', 'Netherlands', 'Australia',
                  'Brazil', 'Colombia', 'Italy', 'Greece']
jobs_df = jobs_df[jobs_df['company_location'].isin(countries)]
jobs_df = jobs_df[jobs_df['employee_residence'].isin(countries)]
jobs_df.sample(10)
```

```
[16]:      work_year      job_title      job_category \
2813      2023      Analytics Engineer      Leadership and Management
856      2023      Machine Learning Engineer      Machine Learning and AI
1076      2023      Data Engineer      Data Engineering
5228      2023      Data Scientist      Data Science and Research
675      2023      Data Scientist      Data Science and Research
5029      2023      Data Scientist      Data Science and Research
2684      2023      Data Analyst      Data Analysis
4363      2023      Data Analyst      Data Analysis
4140      2023      Research Scientist      Data Science and Research
2960      2023      Data Integration Specialist      Data Management and Strategy
```

```
      salary_currency  salary  salary_in_usd  employee_residence \
2813      USD      134000      134000      United States
856      USD      214500      214500      United States
1076      USD      112000      112000      United States
5228      USD      150120      150120      United States
675      USD      92000      92000      United States
5029      USD      85000      85000      United States
2684      USD      130000      130000      United States
4363      USD      113220      113220      United States
4140      USD      220000      220000      United States
2960      USD      85000      85000      Canada
```

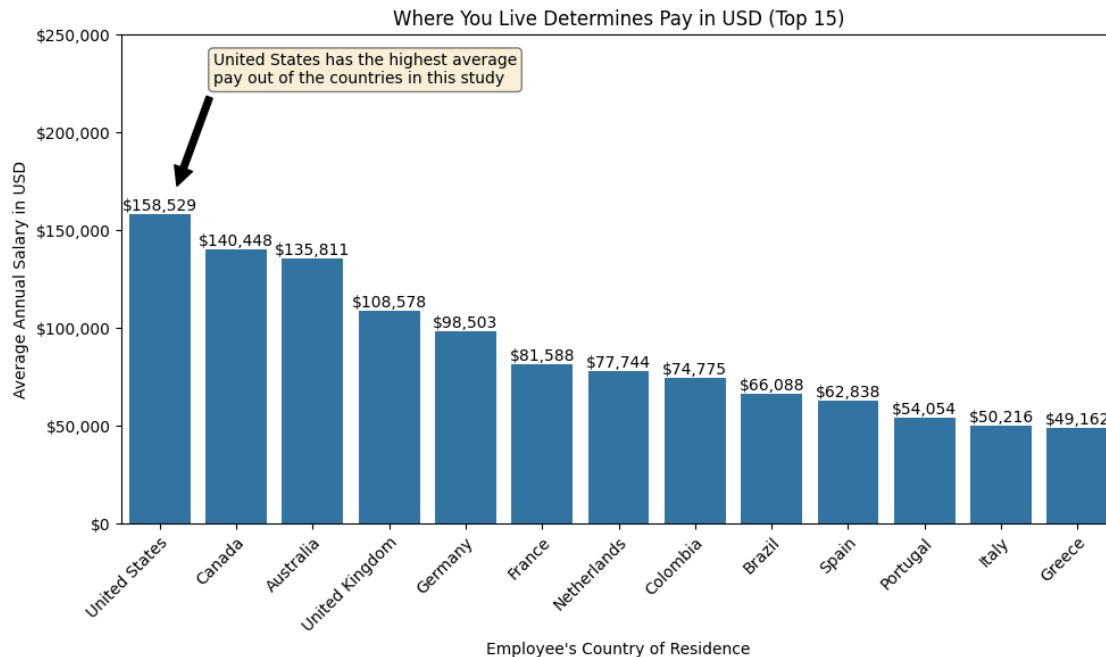
```
      experience_level  employment_type  work_setting  company_location \
2813      Senior      Full-time      Remote      United States
856      Senior      Full-time      In-person      United States
1076      Senior      Full-time      In-person      United States
5228      Senior      Full-time      Remote      United States
675      Mid-level      Full-time      In-person      United States
5029      Mid-level      Full-time      In-person      United States
2684      Senior      Full-time      Remote      United States
4363      Senior      Full-time      In-person      United States
4140      Mid-level      Full-time      Remote      United States
2960      Entry-level      Full-time      In-person      Canada
```

```
company_size
```


2813	M
856	M
1076	M
5228	M
675	M
5029	M
2684	S
4363	M
4140	M
2960	M

0.0.4 Look for interesting items to plot

```
[17]: # Highest pay by employee_residence in USD. Top 15
pay_residence = jobs_df.groupby('employee_residence').agg({'salary_in_usd':
    ↪ 'mean'}). \
    sort_values('salary_in_usd', ascending=False).
    ↪ head(15)
fig, ax = plt.subplots(figsize=(10,6))
props = dict(boxstyle='round', facecolor='wheat', alpha=0.5)
ax.yaxis.set_major_formatter(mtick.StrMethodFormatter('${x:,.0f}'))
ax.set(xlabel='Employee\'s Country of Residence', ylabel='Average Annual Salary_
    ↪ in USD', \
    title='Where You Live Determines Pay in USD (Top 15)', ylim=(0,250000))
sns.barplot(data=pay_residence, x='employee_residence', y='salary_in_usd')
plt.xticks(rotation=45, ha='right', rotation_mode='anchor')
textstr = '\n'.join(('United States has the highest average',
    ↪ 'pay out of the countries in this study'))
ax.annotate(textstr, xy=(0.2,170000), xytext=(0.7,225000), bbox=props,
    ↪ fontsize=10, arrowprops=dict(facecolor='black', shrink=0.05))
ax.bar_label(ax.containers[0], fmt='${:,.0f}')
fig.savefig('../images/highResidenceUSD.png', bbox_inches='tight', dpi=300)
```



```
[18]: # Average salary based on company size
sizes = ('Small', 'Medium', 'Large')
size = jobs_df.groupby('company_size').agg({'salary_in_usd': 'mean'}). \
        sort_values('company_size', ascending=False)
fig, ax = plt.subplots(figsize=(10,6))
sns.barplot(data=size, x='company_size', y='salary_in_usd')
ax.yaxis.set_major_formatter(mtick.StrMethodFormatter('${x:,.0f}'))
x_pos = np.arange(len(sizes))
ax.set_xticks(x_pos, labels=sizes)
ax.set(xlabel='Company Size', ylabel='Average Annual Salary in USD',
       title='Average Annual Salary based on Company Size',
       ylim=(0,180000))
textstr = '\n'.join(('Medium sized organizations',
                     'tend to pay more than smaller',
                     'and larger ones on average'))
props = dict(boxstyle='round', facecolor='wheat', alpha=0.5)
ax.text(0.7, 0.98, textstr, transform=ax.transAxes, fontsize=11,
       verticalalignment='top', bbox=props)
ax.bar_label(ax.containers[0], fmt='${:,.0f}')
fig.savefig('../images/aveCompanySize.png', bbox_inches='tight', dpi=300)
```




```
[24]: # Average salary by job category and company size
fig, ax = plt.subplots(figsize=(12,8))
ax.xaxis.set_major_formatter(mtick.StrMethodFormatter('${x:,.0f}'))
ax.set(xlabel='Annual Average Salary in USD', ylabel='Job Category',
       title='Average Salary by Job Category and Company Size',
       xlim=(0,225000))
hue_order = ['S', 'M', 'L']
bar_colors = ['grey', 'olive', 'purple']
sns.barplot(data=jobs_df, x='salary_in_usd', y='job_category',
            hue='company_size',
            hue_order=hue_order, palette=bar_colors, ci=None)
ax.legend(['Small', 'Medium', 'Large'], title='Company Size')
props = dict(boxstyle='round', facecolor='wheat', alpha=0.5)
text1 = '\n'.join(('As companies grow, they start',
                  'to view data as a resource',
                  'that requires specialized',
                  'planning and governance'))
text2 = '\n'.join(('Medium sized companies pay',
                  'ML/AI engineers more than',
                  'managers. This may indicate',
                  'that ML/AI skills are more',
                  'valued'))
an1 = ax.annotate(text1, xytext=(172000,2.70), xy=(149000,1.4), bbox=props,
                  fontsize=9, arrowprops=dict(facecolor='black', shrink=0.05))
an2 = ax.annotate(text2, xytext=(172000,4.53), xy=(149000,3.2), bbox=props,
```

```

        fontsize=9, arrowprops=dict(facecolor='black', shrink=0.05))
an3 = ax.annotate(text2, xytext=(172000,4.53), xy=(149000,5.0), bbox=props,
        fontsize=9, arrowprops=dict(facecolor='black', shrink=0.05))
fig.savefig('../images/aveCatSize.png', bbox_inches='tight', dpi=300)

```




```

[20]: # Average Salary by job category and work setting
fig, ax = plt.subplots(figsize=(12,8))
ax.xaxis.set_major_formatter(mtick.StrMethodFormatter('${x:,.0f}'))
ax.set(xlabel='Annual Average Salary in USD', ylabel='Job Category',
        title='Average Salary by Job Category and Work Setting in USD',
        xlim=(0,225000))
hue_order = ['In-person', 'Hybrid', 'Remote']
bar_colors = ['lightgrey', 'burlywood', 'royalblue']
props = dict(boxstyle='round', facecolor='wheat', alpha=0.5)
sns.barplot(data=jobs_df, x='salary_in_usd', y='job_category',
        hue='work_setting',
        hue_order=hue_order, palette=bar_colors, ci=None)
ax.legend(title='Work Setting')
text1 = '\n'.join(('No longer do people have to',
        'commute to the office to earn',
        'a good salary'))
ax.text(0.7, 0.55, text1, transform=ax.transAxes, fontsize=9,
        verticalalignment='top', bbox=props)

```

```
fig.savefig('../images/aveCatWork.png', bbox_inches='tight', dpi=300)
```



```
[ ]:
```