

WileyWinters_Week5_Assignment

February 18, 2024

0.0.1 Week 5 Assignment

- Wiley Winters
 - MSDS 670 — Data Visualization
 - 18-FEB-2024
-

0.0.2 Dataset Information

Dataset: Jobs and Salaries in Data Science Metadata: - **work_year**: Year in which data was recorded. - **job_title**: Specific title of the job role. - **job_category**: Classification of the job role into broader categories for easier analysis - **salary_currency**: Currency in which the salary is paid - **salary**: Annual gross salary of the role in the local currency - **salary_in_usd**: Annual gross salary in USD - **employee_residence**: Country of residence - **experience_level**: Classifies the professional experience level of the employee - **employment_type**: Specifies the type of employment such as *full-time*, *part-time*, *contract*, *etc* - **work_setting**: Work setting or environment such as *remote*, *in-person*, or *hybrid* - **company_location**: Country where the company is located - **company_size**: Size of the employer company categorized as *small (S)*, *medium (M)*, and *large (L)*

Formal Reference to Dataset

Qaasim, H. (2023, December). Jobs and Salaries in Data Science. Version 6. Retrieved December 25, 2023 from <https://www.kaggle.com/datasets/hummaamqaasim/jobs-in-data/data>

Import required packages and libraries. Set global configuration items.

```
[1]: import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
import matplotlib.ticker as mtick
from matplotlib import rcParams
import numpy as np

# Suppress Warnings
import warnings
warnings.filterwarnings('ignore')

# Set seaborn style and autoconfig
```

```
sns.set_style('whitegrid')
rcParams.update({'figure.autolayout': True})
```

Read dataset into a Pandas DataFrame

```
[2]: jobs_df = pd.read_csv('data/jobs_in_data.csv')
jobs_df.sample(5)
```

```
[2]:
```

	work_year	job_title	job_category	\
3799	2023	Data Science Manager	Data Science and Research	
54	2023	Research Analyst	Data Science and Research	
819	2023	Business Intelligence Engineer	BI and Visualization	
6483	2023	Data Scientist	Data Science and Research	
2558	2023	Data Scientist	Data Science and Research	

	salary_currency	salary	salary_in_usd	employee_residence	\
3799	USD	245100	245100	United States	
54	USD	64000	64000	United States	
819	USD	127000	127000	Canada	
6483	USD	128750	128750	United States	
2558	USD	100000	100000	United States	

	experience_level	employment_type	work_setting	company_location	\
3799	Senior	Full-time	In-person	United States	
54	Entry-level	Full-time	Remote	United States	
819	Senior	Full-time	In-person	Canada	
6483	Mid-level	Full-time	In-person	United States	
2558	Senior	Full-time	In-person	United States	

	company_size
3799	M
54	M
819	M
6483	M
2558	M

```
[3]: jobs_df.describe().T
```

```
[3]:
```

	count	mean	std	min	25%	\
work_year	9355.0	2022.760449	0.519470	2020.0	2023.0	
salary	9355.0	149927.981293	63608.835387	14000.0	105200.0	
salary_in_usd	9355.0	150299.495564	63177.372024	15000.0	105700.0	

	50%	75%	max
work_year	2023.0	2023.0	2023.0
salary	143860.0	187000.0	450000.0
salary_in_usd	143000.0	186723.0	450000.0

The dataset covers years from 2020 to 2023. In order to not double count some values. I will only work with 2023 data

Check some basic items to see if the dataset requires cleaning or not

```
[4]: print(jobs_df.info())
print('\nNaN Values:\n', jobs_df.isna().sum())
print('\nDuplicates: ', jobs_df.duplicated().sum())
print('\nSize: ', jobs_df.size)
print('\nDistribution:\n', jobs_df.describe().T)
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 9355 entries, 0 to 9354
Data columns (total 12 columns):
#   Column                Non-Null Count  Dtype
---  -
0   work_year              9355 non-null   int64
1   job_title              9355 non-null   object
2   job_category           9355 non-null   object
3   salary_currency        9355 non-null   object
4   salary                 9355 non-null   int64
5   salary_in_usd          9355 non-null   int64
6   employee_residence     9355 non-null   object
7   experience_level        9355 non-null   object
8   employment_type        9355 non-null   object
9   work_setting           9355 non-null   object
10  company_location       9355 non-null   object
11  company_size           9355 non-null   object
dtypes: int64(3), object(9)
memory usage: 877.2+ KB
None
```

```
NaN Values:
work_year      0
job_title      0
job_category   0
salary_currency 0
salary         0
salary_in_usd  0
employee_residence 0
experience_level 0
employment_type 0
work_setting   0
company_location 0
company_size   0
dtype: int64
```

```
Duplicates: 4014
```

Size: 112260

Distribution:

	count	mean	std	min	25%	50%	75%	max
work_year	9355.0	2022.760449	0.519470	2020.0	2023.0	2023.0	2023.0	2023.0
salary	9355.0	149927.981293	63608.835387	14000.0	105200.0	143860.0	187000.0	450000.0
salary_in_usd	9355.0	150299.495564	63177.372024	15000.0	105700.0	143000.0	186723.0	450000.0


Looks like there is a lot of duplicates. I will remove them.

```
[5]: jobs_df.drop_duplicates(keep='first', inplace=True)
jobs_df.duplicated().sum()
```


[5]: 0

0.0.3 Look for interesting items to plot


```
[6]: salary_years = jobs_df.groupby('work_year').agg({'salary_in_usd': 'mean'}). \
      sort_values('work_year')
fig, ax = plt.subplots(figsize=(10,6))
fmt = '${x:,.0f}'
tick = mtick.StrMethodFormatter(fmt)
ax.yaxis.set_major_formatter(tick)
ax.set(xlabel='Year',ylabel='Average Salary in USD',title='Average Salary over_
↳the Years')
sns.barplot(data=salary_years, y='salary_in_usd', x='work_year')
fig.savefig('images/yearAveSalary.png', bbox_inches='tight', dpi=300)
```




```
[7]: fig, ax = plt.subplots(figsize=(10,6))
ax.set(ylabel='Job Category',title='Job Count by Category',xlabel='Count')
sns.countplot(ax=ax, data=jobs_df, y='job_category',
              order=jobs_df['job_category']. \
                    value_counts().index)
ax.bar_label(ax.containers[0])
fig.savefig('images/jobCatCount.png', bbox_inches='tight', dpi=300)
```



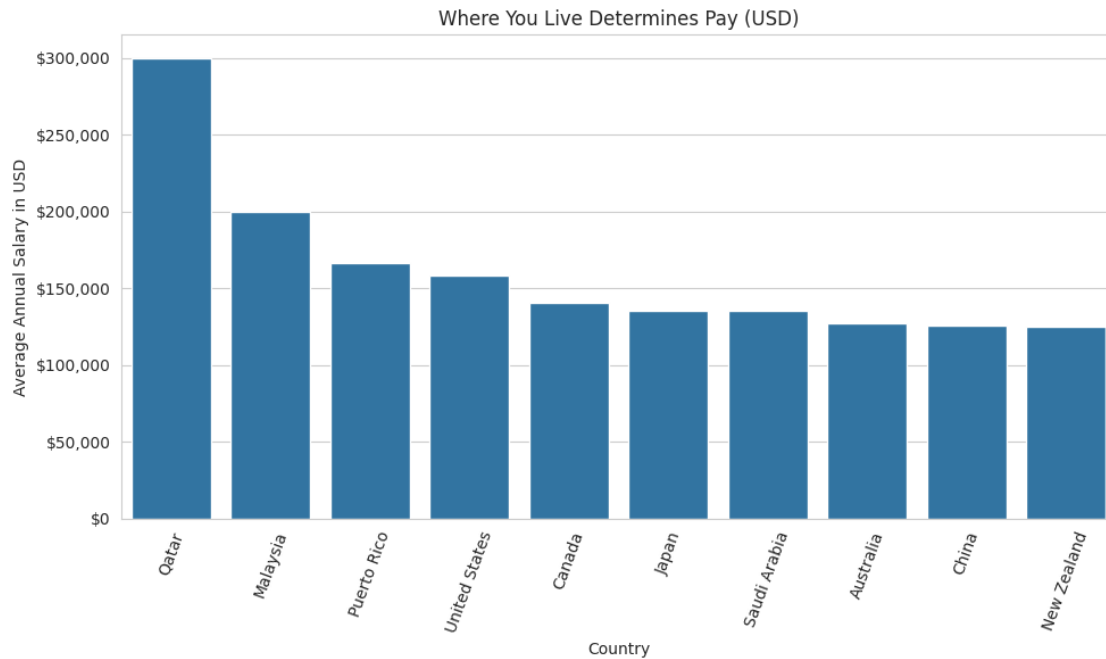
```
[8]: # Highest paying categories in USD
salary = jobs_df.groupby('job_category').agg({'salary_in_usd': 'mean'}). \
        sort_values('salary_in_usd', ascending=False)
fig, ax = plt.subplots(figsize=(10,6))
fmt = '${x:,.0f}'
tick = mtick.StrMethodFormatter(fmt)
ax.xaxis.set_major_formatter(tick)
ax.set(xlabel='Annual Average Salary (USD)',ylabel='Job Category', \
        title='Highest Paying Job Categories (USD)')
sns.barplot(data=salary, y='job_category', x='salary_in_usd')
fig.savefig('images/highJobCatUSD.png', bbox_inches='tight', dpi=300)
```



```
[9]: # Highest pay by company_location in USD. Top 10
pay_country = jobs_df.groupby('company_location').agg({'salary_in_usd':
    ↪ 'mean'}). \
    sort_values('salary_in_usd', ascending=False).
    ↪ head(10)
fig, ax = plt.subplots(figsize=(10,6))
fmt = '${x:,.0f}'
tick = mtick.StrMethodFormatter(fmt)
ax.yaxis.set_major_formatter(tick)
ax.set(xlabel='Country',ylabel='Average Annual Salary in USD', \
    title='Highest Paying Countries (USD)')
sns.barplot(data=pay_country, x='company_location', y='salary_in_usd')
plt.xticks(rotation=70)
fig.savefig('images/highCountryUSD.png', bbox_inches='tight', dpi=300)
```



```
[10]: # Highest pay by employee_residence in USD. Top 10
pay_residence = jobs_df.groupby('employee_residence').agg({'salary_in_usd':
    ↪ 'mean'}). \
    sort_values('salary_in_usd', ascending=False).
    ↪ head(10)
fig, ax = plt.subplots(figsize=(10,6))
fmt = '${x:,.0f}'
tick = mtick.StrMethodFormatter(fmt)
ax.yaxis.set_major_formatter(tick)
ax.set(xlabel='Country',ylabel='Average Annual Salary in USD', \
    title='Where You Live Determines Pay (USD)')
sns.barplot(data=pay_residence, x='employee_residence', y='salary_in_usd')
plt.xticks(rotation=70)
fig.savefig('images/highResidenceUSD.png', bbox_inches='tight', dpi=300)
```







```
[11]: # Average Salary by company_size
size = jobs_df.groupby('company_size').agg({'salary_in_usd': 'mean'}). \
        sort_values('company_size', ascending=False)
fig, ax = plt.subplots(figsize=(10,6))
fmt = '${x:,.0f}'
tick = mtick.StrMethodFormatter(fmt)
ax.yaxis.set_major_formatter(tick)
ax.set(xlabel='Company Size', ylabel='Average Annual Salary in USD', \
        title='Average Annual Salary based on Company Size (USD)')
sns.barplot(data=size, x='company_size', y='salary_in_usd')
fig.savefig('images/aveCompanySize.png', bbox_inches='tight', dpi=300)
```



```
[12]: sizes = {'S': 'Small', 'M': 'Medium', 'L': 'Large'}

for key in sizes:
    size = jobs_df[jobs_df['company_size'] == key]
    comp_size = size.groupby('job_category').agg({'salary_in_usd': 'mean'}). \
        sort_values('salary_in_usd', ascending=False)
    fig, ax = plt.subplots(figsize=(10,6))
    ax.yaxis.set_major_formatter(mtick.StrMethodFormatter('${x:,.0f}'))
    ax.set(xlabel='Company Size', ylabel='Average Annual Salary in USD', \
        title=f'Highest Paying Job Categories by Company Size (USD): \
↪{sizes[key]}')
    plt.xticks(rotation=70)
    sns.barplot(data=comp_size, x='job_category', y='salary_in_usd')
    fig.savefig('images/'+sizes[key]+'Cat.png', bbox_inches='tight', dpi=300)
```



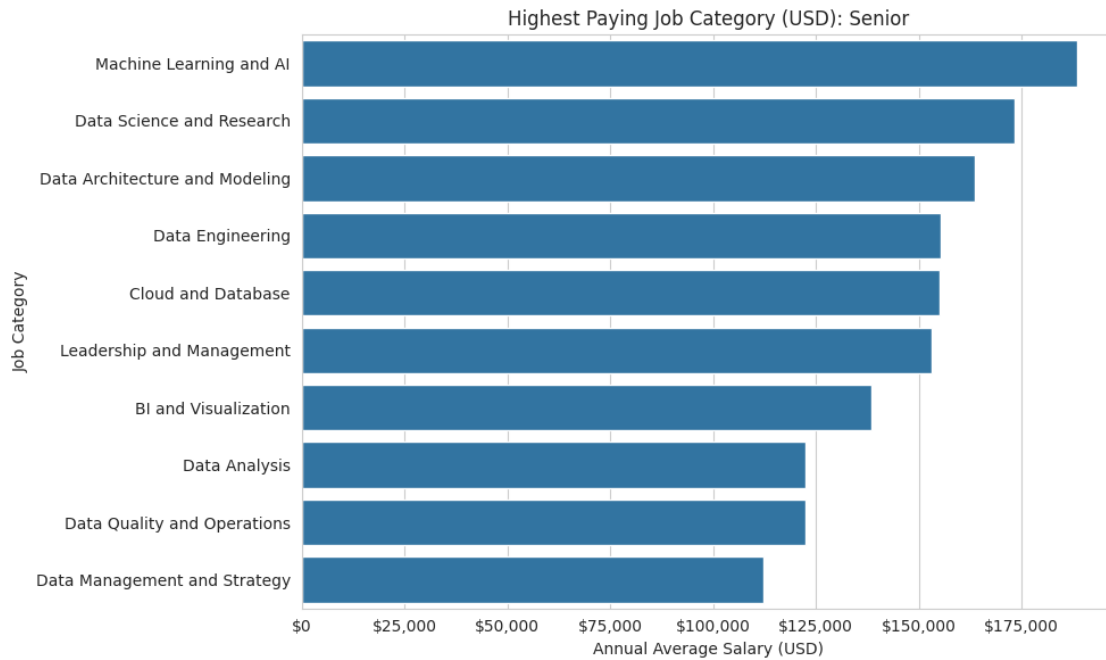



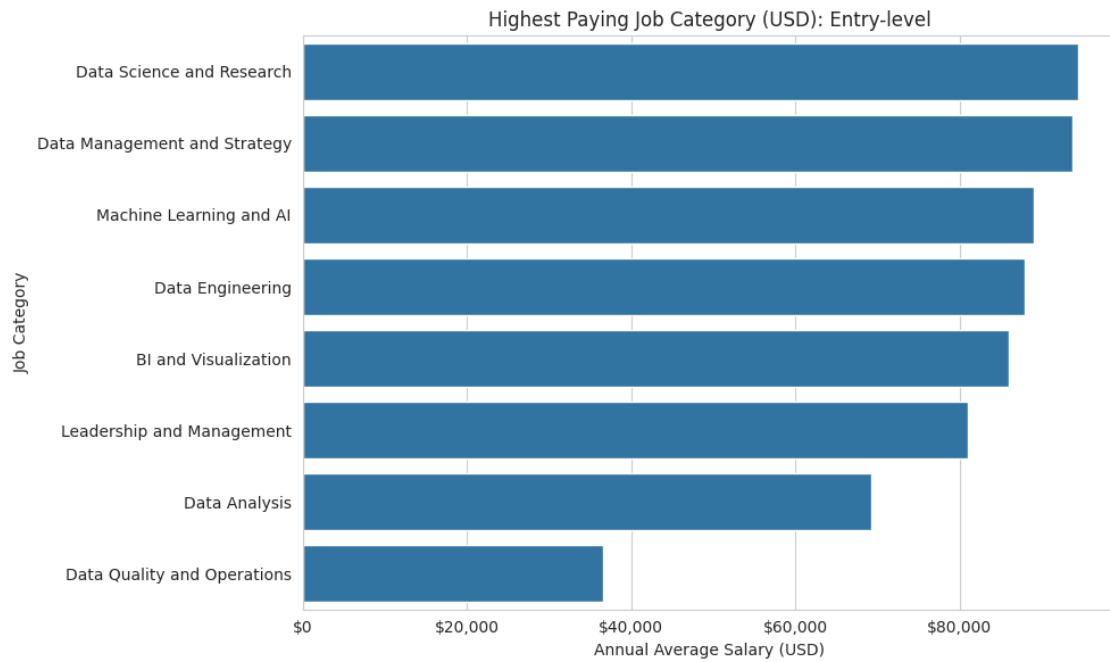
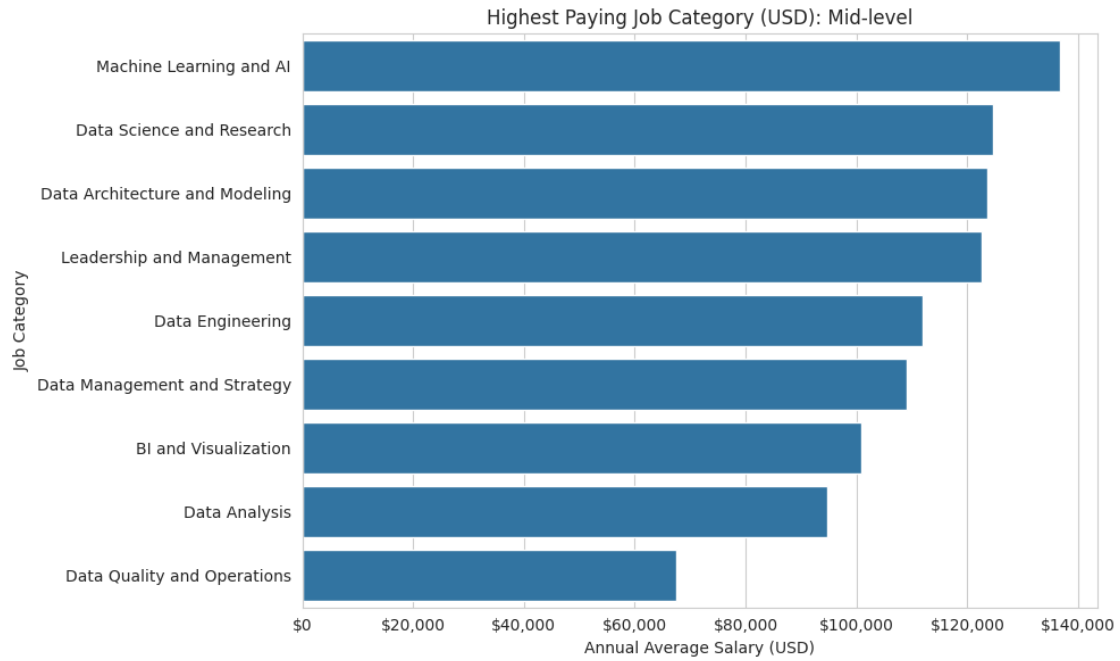
```
[13]: # Salary based on experience level
experience = jobs_df.groupby('experience_level').agg({'salary_in_usd': 'mean'}).
    \
        sort_values('salary_in_usd', ascending=False)
fig, ax = plt.subplots(figsize=(10,6))
fmt = '${x:,.0f}'
tick = mtick.StrMethodFormatter(fmt)
ax.yaxis.set_major_formatter(tick)
ax.set(xlabel='Experience Level', ylabel='Average Annual Salary in USD', \
       title='Average Salary by Experience Level (USD)')
sns.barplot(data=experience, x='experience_level', y='salary_in_usd')
fig.savefig('images/salaryByexperience.png', bbox_inches='tight', dpi=300)
```



Explore each experience level to see what job categories pay the most in USD

```
[14]: # Look at salary based on experience level for all levels
experiences = ['Executive', 'Senior', 'Mid-level', 'Entry-level']
for i in range(len(experiences)):
    exp = jobs_df[jobs_df['experience_level'] == experiences[i]]
    salary = exp.groupby('job_category').agg({'salary_in_usd': 'mean'}). \
        sort_values('salary_in_usd', ascending=False)
    fig, ax = plt.subplots(figsize=(10,6))
    ax.xaxis.set_major_formatter(mtick.StrMethodFormatter('${x:,.0f}'))
    ax.set(xlabel='Annual Average Salary (USD)', ylabel='Job Category', \
        title=f'Highest Paying Job Category (USD): {experiences[i]}')
    sns.barplot(data=salary, y='job_category', x='salary_in_usd')
    fig.savefig('images/'+experiences[i]+'.png', bbox_inches='tight', dpi=300)
```





[]: